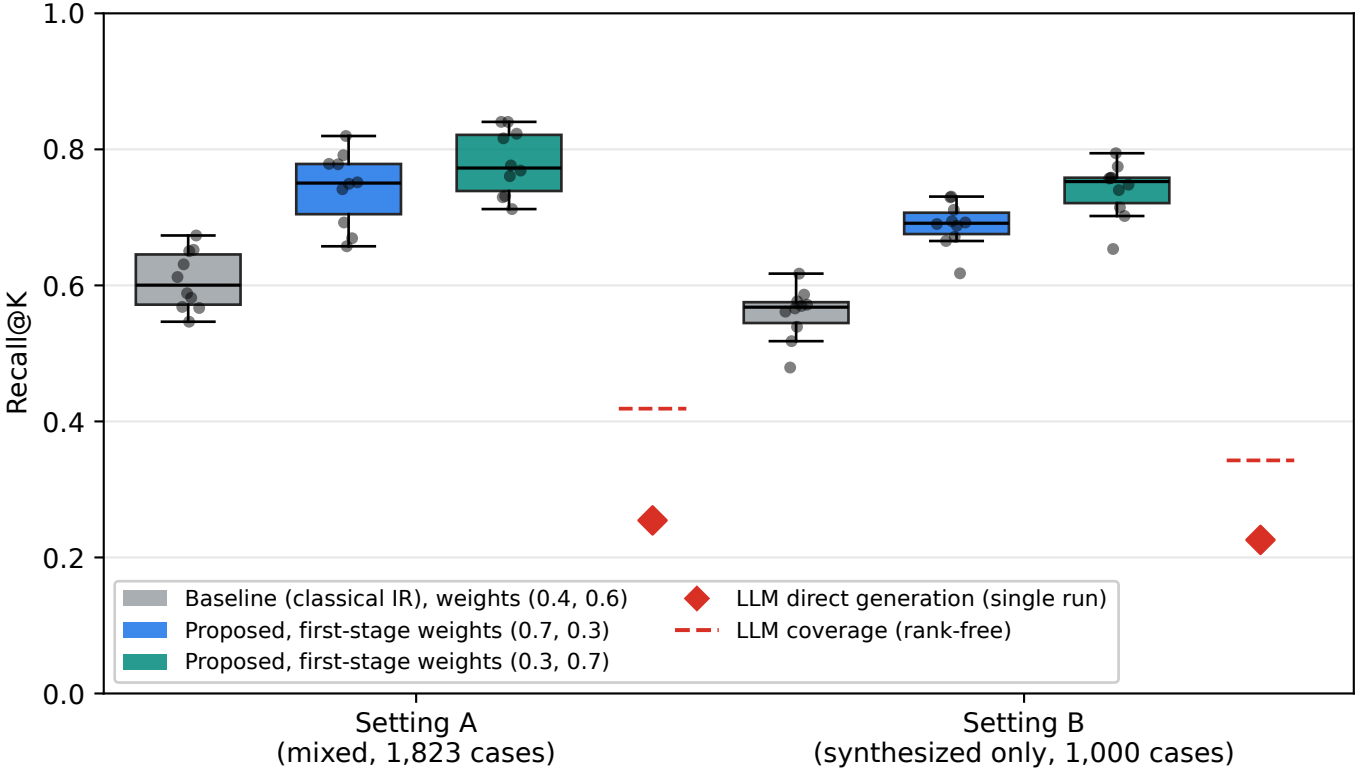
1.0
0.8
0.6
0.4
0.2
0.0
Recall@K
Baseline (classical IR), weights (0.4, 0.6)
Proposed, first-stage weights (0.7, 0.3)
Proposed, first-stage weights (0.3, 0.7)
LLM direct generation (single run)
LLM coverage (rank-free)
Setting A
(mixed, 1,823 cases)
Setting B
(synthesized only, 1,000 cases)